

Food web reconstruction through phylogenetic transfer of low-rank network representation

Tanya Strydom^{1,2,‡} Salomé Bouskila^{1,‡} Francis Banville^{1,3,2} Ceres Barros⁴ Dominique Caron^{5,2}
Maxwell J Farrell⁶ Marie-Josée Fortin⁶ Victoria Hemming⁴ Benjamin Mercier^{3,2} Laura
J. Pollock^{5,2} Rogini Runghen⁷ Giulio V. Dalla Riva⁸ Timothée Poisot^{1,2}

¹ Département de Sciences Biologiques, Université de Montréal, Montréal, Canada ² Québec Centre for Biodiversity Sciences, Montréal, Canada ³ Université de Sherbrooke, Sherbrooke, Canada

⁴ Department of Forest Resources Management, University of British Columbia, Vancouver, Canada

⁵ Department of Biology, McGill University, Montréal, Canada ⁶ Department of Ecology & Evolutionary Biology, University of Toronto, Toronto, Canada ⁷ Centre for Integrative Ecology, School of Biological Sciences, University of Canterbury, Canterbury, New Zealand ⁸ School of Mathematics and Statistics, University of Canterbury, Canterbury, New Zealand

‡ These authors contributed equally to the work

Correspondance to:

Timothée Poisot — timothee.poisot@umontreal.ca

Despite their importance in many ecological processes, collecting data and information on ecological interactions is an exceedingly challenging task. For this reason, large parts of the world have a data deficit when it comes to species interactions, and how the resulting networks are structured. As data collection alone is unlikely to be sufficient, community ecologists must adopt predictive methods. Here we develop such a method, relying on graph embedding and transfer learning to assemble a predicted list of trophic interactions between Canadian mammals. This interaction list is derived from the European food web, despite sharing only 5% of common species with Canada. The results of the predictive model are compared against databases of recorded pairwise interactions, showing that we correctly recover over 95% of known interactions. We provide guidance on how this method can be adapted by substituting some approaches or predictors in order to make it more generally applicable.

1 Introduction

2 There are two core challenges we are faced with in furthering our understanding of ecological networks
3 across space, particularly at macro-ecologically relevant scales (e.g. Trøjelsgaard & Olesen 2016). First,
4 networks within a location are difficult to sample properly (Jordano 2016a, b), resulting in a widespread
5 “Eltonian shortfall” (Hortal *et al.* 2015), *i.e.* a lack of knowledge about inter and intra specific
6 relationships. This first challenge has been, in large part, addressed by the recent emergence of a suite of
7 methods aiming to predict interactions within *existing* networks, many of which are reviewed in Strydom
8 *et al.* (2021a). Second, recent analyses based on collected data (Poisot *et al.* 2021a) or metadata (Cameron
9 *et al.* 2019) highlight that ecological networks are currently studied in a biased subset of space and
10 bioclimates, which impedes our ability to generalize any local understanding of network structure.
11 Meaning that, although the framework to address incompleteness *within* networks exists, there would still
12 be regions for which, due to a *lack* of local interaction data, we are unable to infer potential species
13 interactions. Having a general solution for inferring **plausible interactions** (despite the unavailability of
14 interaction data) could be the catalyst for significant breakthroughs in our ability to start thinking about
15 species interaction networks over large spatial scales. **In a recent overview of the field of ecological**
16 **network prediction, Strydom *et al.* (2021a) identified two challenges of interest to the prediction**
17 **of interactions at large scales. First, there is a relative scarcity of relevant data in most places**
18 **globally – paradoxically, this restricts our ability to infer interactions to locations where**
19 **inference is perhaps the least required; second, accurate predictions often demand accurate**
20 **predictors, and the lack of methods that can leverage small amount of data is a serious**
21 **impediment to our predictive ability globally.**

22 Here, we present a general method **to recommend possible trophic interactions, relying on** the
23 transfer learning of network representations, **specifically by using** similarities of species in a
24 biologically/ecologically relevant proxy space (e.g. shared morphology or ancestry). Transfer learning is a
25 machine learning methodology that uses the knowledge gained from solving one problem and applying it
26 to a related (destination) problem (Pan & Yang 2010; Torrey & Shavlik 2010). In this instance, we solve the
27 problem of predicting trophic interactions between species, based on knowledge extracted from another
28 species pool for which interactions are known by using phylogenetic structure as a medium for transfer.

29 There is a plurality of measures of species similarities that can be used for metaweb reconstruction (see e.g.

Morales-Castilla *et al.* 2015); however, phylogenetic proximity has several desirable properties when working at large scales. Gerhold *et al.* (2015) made the point that phylogenetic signal captures diversification of characters (large macro-evolutionary process), but not necessarily community assembly (fine ecological process); Dormann *et al.* (2010) previously found very similar conclusions. Interactions tend to reflect a phylogenetic signal because they have a conserved pattern of evolutionary convergence that encompasses a wide range of ecological and evolutionary mechanisms (Cavender-Bares *et al.* 2009; Mouquet *et al.* 2012), and - most importantly - retain this signal even when it is not detectable at the community scale (Hutchinson *et al.* 2017; Poisot & Stouffer 2018). Finally, species interactions at macro-ecological scales seem to respond mostly to macro-evolutionary processes (Price 2003); which is evidenced by the presence of conserved backbones in food webs (Dalla Riva & Stouffer 2016; Mora *et al.* 2018), strong evolutionary signature on prey choice (Stouffer *et al.* 2012), and strong phylogenetic signature in food web intervality (Eklöf & Stouffer 2016). Phylogenetic reconstruction has also previously been used within the context of ecological networks, namely understanding ancestral plant-insect interactions (Braga *et al.* 2021). Taken together, these considerations suggest that phylogenies can reliably be used to transfer knowledge on species interactions.

[Figure 1 about here.]

Our methodology is outlined in fig. 1, where we provide an illustration based on learning the embedding of a metaweb of trophic interactions for European mammals (known interactions; Maiorano *et al.* 2020b, a) and, based on phylogenetic relationships between mammals globally (*i.e.*, phylogenetic tree Upham *et al.* 2019), infer a metaweb for the Canadian mammalian species pool (interactions are treated as unknown in this instance). Following the definition of Dunne (2006), a metaweb is a network analogue to the concept of a regional species pool. Specifically, a metaweb is an inventory of all *possible* interactions within species likely to occur within a spatially delimited area (the network γ -diversity, in a sense). The metaweb is, therefore, *not* a prediction of the food web at any specific locale within the frontiers of the spatial area it recovers, and will in fact have a different structure (notably by having a larger connectance; see *e.g.* Wood *et al.* 2015). These local food webs are expected to be a subset of both the species and the interactions of their metaweb, and have been called “metaweb realizations” (Poisot *et al.* 2015). The difference between a food web at a specific location and the metaweb

has to do with a variety of mechanisms, including species co-occurrence, local environmental conditions, and local distribution of functional traits. Nevertheless, the metaweb represents the total of functional, phylogenetic, and macroecological processes (Morales-Castilla *et al.* 2015), and therefore still holds valuable ecological information. Because the metaweb can be down-sampled given appropriate knowledge of local species composition (the network α -diversity, in a sense), it is possible to infer what may drive the structure of food webs at finer spatial scales. This has been done for example for tree-gallers-parasitoid systems (Gravel *et al.* 2018), fish trophic interactions (Albouy *et al.* 2019), tetrapods trophic interactions (O'Connor *et al.* 2020), and crop-pests networks (Grünig *et al.* 2020). Whereas the original metaweb definition was based on presence/absence, we focus on *probabilistic* metawebs; not only does our method recommend interactions that may exist, it gives each interaction a score that is mathematically equivalent to the chance of success of a Bernoulli trial (see *e.g.* Poisot *et al.* 2016), which allows properly weigh interactions as a function of how likely they are.

Our case study shows that phylogenetic transfer learning is an effective approach to **the generation of probabilistic metawebs**. This showcases that although the components (species) that make up the Canadian and European communities may be *minimally* shared (**the overall species overlap is less than 4%**), if the medium (proxy space) selected in the transfer step is biologically plausible, we can still effectively learn from the known network and make biologically relevant predictions of interactions. **Indeed, as we detail in the result, when validated against known but fractional data of trophic interactions between Canadian mammals, our model achieves a predictive accuracy of approx. 91%.** It should be reiterated that the framework presented in fig. 1 is amenable to changes; notably, the measure of similarity may not be phylogeny, and can be replaced by information on foraging (Beckerman *et al.* 2006), cell-level mechanisms (Boeckaerts *et al.* 2021), or a combination of traits and phylogenetic structure (Stock 2021). **Most importantly, although we focus on a trophic system, it is an established fact that different types of interactions do themselves interact (Kéfi *et al.* 2012; see *e.g.* Kawatsu *et al.* 2021). Future development of metaweb inference techniques should cover the prediction of multiple interaction types.**

86 Data used for the case study

87 We use data from the European metaweb assembled by Maiorano *et al.* (2020b).

88 TANYA TODO description of the data

89 **We induced** the subgraph corresponding to all mammals by matching species names in the original
90 network to the GBIF taxonomic backbone (GBIF Secretariat 2021) and retaining all those who matched to
91 mammals. This serves a dual purpose 1) to extract only mammals from the European network and 2) to
92 match and standardize species names when aggregating the different data sources further downstream
93 (which is an important consideration when combining datasets (Grenié *et al.* 2021)). All nodes had valid
94 matches to GBIF at this step, and so this backbone is used for all name reconciliation steps as outlined
95 below.

96 The European metaweb represents the knowledge we want to learn and transfer; the phylogenetic
97 similarity of mammals here represents the information for transfer. We used the mammalian consensus
98 supertree by Upham *et al.* (2019), for which all approximately 6000 names have been similarly matched
99 to their GBIF valid names. This step allows us to place each node of the mammalian European metaweb
100 in the phylogeny.

101 The destination problem to which we want to transfer knowledge is the trophic interactions between
102 mammals in Canada. We obtained the list of extant species from the IUCN checklist, and selected the
103 terrestrial and semi-aquatic species (this corresponds to the same selection that was applied by Maiorano
104 *et al.* (2020b) in the European metaweb). The IUCN names were, as previously, reconciled against GBIF to
105 have an exact match to the taxonomy.

106 After taxonomic cleaning and reconciliation as outlined in the following sections, the mammalian
107 European metaweb has 260 species, and the Canadian species pool has 163; of these, 17 (about 4% of the
108 total) are shared, and 89 species from Canada (54%) had at least one congeneric species in Europe. The
109 similarity for both species pools predictably increases with higher taxonomic order, with 19% of shared
110 genera, 47% of shared families, and 75% of shared orders; for the last point, Canada and Europe each had a
111 single unique order (*Didelphimorphia* for Canada, *Erinaceomorpha* for Europe).

112 In the following sections, we describe the representational learning step applied to European data, the
113 transfer step through phylogenetic similarity, and the generation of a probabilistic metaweb for the

114 destination species pool.

115 **Method description**

116 The **core point of our** method is the transfer of knowledge of a known **ecological** network, in order to
117 predict interactions between species from another location **at which the network is unknown (or**
118 **partially known)**. In fig. 1, we give a high-level overview of the approach; in the example around which
119 this manuscript is built (leveraging detailed knowledge about binary trophic interactions between
120 Mammalia in Europe to predict the less known trophic interactions between closely phylogenetically
121 related Mammalia in Canada), we use a series of specific steps for network embedding, trait inference,
122 network prediction and thresholding.

123 Specifically, our approach can be summarized as follows: from the known network in Europe, we use a
124 truncated Singular Value Decomposition (t-SVD; Halko *et al.* 2011) to generate latent traits representing a
125 low-dimensional embedding of the network; these traits give an unbiased estimate of the node's position
126 in the latent feature spaces. Then, we map these latent traits onto a reference phylogeny (other
127 distance-based measures of species proximity that allow for the inference of features in the latent space
128 can be used, for example the dissimilarity in functional traits). Based on the reconstructed latent traits for
129 species in the destination species pool, a Random Dot Product Graph model (hereafter RDPG; Young &
130 Scheinerman 2007) predicts the interaction between species through a function of the nodes' features
131 through matrix multiplication. Thus, from latent traits and node position, we can infer interactions.

132 **The method we develop is, ecologically speaking, a “black box,” i.e. an algorithm that can be**
133 **understood mathematically, but whose component parts are not always directly tied to**
134 **ecological processes. There is a growing realization in machine learning that (unintentional)**
135 **black box algorithms are not necessarily a bad thing Holm (2019), as long as their constituent**
136 **parts can be examined (which is the case with our method). But more importantly, data hold**
137 **more information than we may thought; as such, even algorithms that are disconnected from**
138 **the model can make correct guesses most of the time (Halevy *et al.* 2009); in fact, in an instance**
139 **of ecological forecasting of spatio-temporal systems, model-free approaches (i.e. drawing all of**
140 **their information from the data) outperformed model-informed ones (Perretti *et al.* 2013).**

Implementation and code availability

The entire pipeline is implemented in *Julia* 1.6 (Bezanson *et al.* 2017) and is available under the permissive MIT License at <https://osf.io/2zwqm/>. The taxonomic cleanup steps are done using `GBIF.jl` (Dansereau & Poisot 2021). The network embedding and analysis is done using `EcologicalNetworks.jl` (Poisot *et al.* 2019; Banville *et al.* 2021). The phylogenetic simulations are done using `PhyloNetworks.jl` (Solís-Lemus *et al.* 2017) and `Phylo.jl` (Reeve *et al.* 2016). A complete `Project.toml` file specifying the full tree of dependencies is available alongside the code. This material also includes a fully annotated copy of the entire code required to run this project (describing both the intent of the code and discussing some technical implementation details), a vignette for every step of the process, and a series of Jupyter notebooks with the text and code. The pipeline can be executed on a laptop in a matter of minutes, and therefore does not require extensive computational power.

Step 1: Learning the origin network representation

The first step in transfer learning is to learn the structure of the original dataset. In order to do so, we rely on an approach inspired from representational learning, where we learn a *representation* of the metaweb (in the form of the latent subspaces), rather than a list of interactions (species *a* eats *b*). This approach is conceptually different from other metaweb-scale predictions (e.g. Albouy *et al.* 2019), in that the metaweb representation is easily transferable. Specifically, we use RDPG to create a number of latent variables that can be combined into an approximation of the network adjacency matrix. RDPG results are known to have strong phylogenetic signal, and to capture the evolutionary backbone of food webs (Dalla Riva & Stouffer 2016); **in other words, the latent variables of an RDPG can be mapped onto a phylogenetic tree, and phylogenetically close predators should share phylogenetically close preys.** In addition, recent advances show that the latent variables produced this way can be used to predict *de novo* network edges. **Interestingly, the latent variables do not need to be produced by decomposing the network itself; in a recent contribution, Runghen *et al.* (2021) show that deep artificial neural networks are able to reconstruct the left and right subspaces of an RDPG, in order to predict human movement networks from individual/location metadata. This is an exciting opportunity, as it opens up the possibility of using additional predictors.**

The latent variables are created by performing a truncated Singular Value Decomposition (t-SVD) on the

adjacency matrix. SVD is an appropriate embedding of ecological networks, which has recently been shown to both capture their complex, emerging properties (Strydom *et al.* 2021b) and to allow highly accurate prediction of the interactions within a single network (Poisot *et al.* 2021b). Under SVD, an adjacency matrix \mathbf{A} (where $\mathbf{A}_{m,n} \in \mathbb{B}$ where 1 indicates predation and 0 an absence thereof) is decomposed into three components resulting in $\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}'$. Here, $\mathbf{\Sigma}$ is a $m \times n$ diagonal matrix and contains only singular (σ) values along its diagonal, \mathbf{U} is a $m \times m$ unitary matrix, and \mathbf{V}' a $n \times n$ unitary matrix. Truncating the SVD removes additional noise in the dataset by omitting non-zero and/or smaller σ values from $\mathbf{\Sigma}$ using the rank of the matrix. Under a t-SVD $\mathbf{A}_{m,n}$ is decomposed so that $\mathbf{\Sigma}$ is a square $r \times r$ diagonal matrix (whith $1 \leq r \leq r_{full}$ where r_{full} is the full rank of \mathbf{A} and r the rank at which we truncate the matrix) containing only non-zero σ values. Additionally, \mathbf{U} is now a $m \times r$ semi unitary matrix and \mathbf{V}' a $n \times r$ semi-unitary matrix. **As an aside, most ecologists are indirectly familiar with SVD: Principal Component Analysis is a special case of SVD, which is more sensitive to numerical instabilities (see notably Shlens 2014).**

The specific rank at which the SVD ought to be truncated is a difficult question. The purpose of SVD is to remove the noise (expressed at high dimensions) and to focus on the signal, (expressed at low dimensions). In datasets with a clear signal/noise demarcation, a scree plot of $\mathbf{\Sigma}$ can show a sharp drop at the rank where noise starts (Zhu & Ghodsi 2006). Because the European metaweb is almost entirely known, the amount of noise (uncertainty) is low; this is reflected in fig. 2 (left), where the scree plot shows no important drop, and in fig. 2 (right) where the proportion of variance explained increases smoothly at higher dimensions. For this reason, we default back to a threshold that explains 60% of the variance in the underlying data, corresponding to 12 dimensions - *i.e.* a tradeoff between accuracy and a reduced number of features.

An RDPG estimates the probability of observing interactions between nodes (species) as a function of the nodes' latent variables, **and is a way to turn a SVD (which decompose a matrix into three) into two matrices that can be multiplied to provide an approximation of the network.** The latent variables used for the RDPG, called the left and right subspaces, are defined as $\mathcal{L} = \mathbf{U}\sqrt{\mathbf{\Sigma}}$, **and** $\mathcal{R} = \sqrt{\mathbf{\Sigma}}\mathbf{V}'$ – using the full rank of \mathbf{A} , $\mathcal{L}\mathcal{R} = \mathbf{A}$, and using any smaller rank results in $\mathcal{L}\mathcal{R} \approx \mathbf{A}$. Using a rank of 1 for the t-SVD provides a first-order approximation of the network. **One advantage of using a RDPG rather than a SVD is that the number of components to estimate decreases; notably, one does not have to estimate the eigenvalues of the SVD. Furthermore, the two subspaces can be directly multiplied to yield a network.**

Because RDPG relies on matrix multiplication, the higher dimensions essentially serve to make specific interactions converge towards 0 or 1; therefore, for reasonably low ranks, there is no guarantee that the values in the reconstructed network will be within the unit range. In order to determine what constitutes an appropriate threshold for probability, we performed the RDPG approach on the European metaweb, and evaluated the probability threshold by treating this as a binary classification problem, specifically assuming that both 0 and 1 in the European metaweb are all true. Given the methodological details given in Maiorano *et al.* (2020b) and O'Connor *et al.* (2020), this seems like a reasonable assumption, although one that does not hold for all metawebs. We used the thresholding approach presented in Poisot *et al.* (2021b), and picked a cutoff that maximized Youden's J statistic (a measure of the informedness (trust) of predictions; Youden (1950)); the resulting cutoff was 0.22, and gave an accuracy above 0.99. **In Supp. Mat. 1, we provide several lines of evidence that using the entire network to estimate the threshold does not lead to overfitting; that using a subset of species would yield the same threshold; that decreasing the quality of the original data by adding or removing interactions would minimally affect the predictive accuracy of RDPG applied to the European metaweb; and that the networks reconstructed from artificially modified data are reconstructed with the correct ecological properties.**

The left and right subspaces for the European metaweb, accompanied by the threshold for prediction, represent the knowledge we seek to transfer. In the next section, we explain how we rely on phylogenetic similarity to do so.

Steps 2 and 3: Transfer learning through phylogenetic relatedness

In order to transfer the knowledge from the European metaweb to the Canadian species pool, we performed ancestral character estimation using a Brownian motion model, which is a conservative approach in the absence of strong hypotheses about the nature of phylogenetic signal in the network decomposition (Litsios & Salamin 2012). This uses the estimated feature vectors for the European mammals to create a state reconstruction for all species (conceptually something akin to a trait-based mammalian phylogeny using generality and vulnerability traits) and allows us to impute the missing (latent) trait data for the Canadian species that are not already in the European network; as we are focused

on predicting contemporary interactions, we only retained the values for the tips of the tree. We assumed that all traits (*i.e.* the feature vectors for the left and right subspaces) were independent, which is a reasonable assumption as every trait/dimension added to the t-SVD has an *additive* effect to the one before it. Note that the Upham *et al.* (2019) tree itself has some uncertainty associated to inner nodes of the phylogeny. In this case study, we have decided to not propagate this uncertainty, as it would complexify the process. The Brownian motion algorithm returns the *average* value of the trait, and its upper and lower bounds. Because we do not estimate other parameters of the traits' distributions, we considered that every species trait is represented as a uniform distribution between these bounds. **The choice of the uniform distribution was made because the algorithm returns a minimum and maximum point estimate for the value, and given this information, the uniform distribution is the one with maximum entropy. Had all mean parameters estimates been positive, the exponential distribution would have been an alternative, but this is not the case for the subspaces of an RDPG. In order to examine the consequences of the choice of distribution, we estimated the variance per latent variable per node to use a normal distribution; as we show in Supp. Mat. 2, this decision results in dramatically over-estimating the number and probability of interactions, and therefore we keep the discussions in the main text to the uniform case.** The inferred left and right sub-spaces for the Canadian species pool ($\hat{\mathcal{L}}$ and $\hat{\mathcal{R}}$) have entries that are distributions, representing the range of values for a given species at a given dimension.

These objects represent the transferred knowledge, which we can use for prediction of the Canadian metaweb.

Step 4: Probabilistic prediction of the destination network

The phylogenetic reconstruction of $\hat{\mathcal{L}}$ and $\hat{\mathcal{R}}$ has an associated uncertainty, represented by the breadth of the uniform distribution associated to each of their entries. Therefore, we can use this information to assemble a *probabilistic* metaweb in the sense of Poisot *et al.* (2016), *i.e.* in which every interaction is represented as a single, independent, Bernoulli event of probability p .

[Figure 3 about here.]

Specifically, we have adopted the following approach. For every entry in $\hat{\mathcal{L}}$ and $\hat{\mathcal{R}}$, we draw a value from its distribution. This results in one instance of the possible left ($\hat{\ell}$) and right (\hat{r}) subspaces for the

255 Canadian metaweb. These can be multiplied, to produce one matrix of real values. Because the entries in
256 $\hat{\mathcal{L}}$ and $\hat{\mathcal{R}}$ are in the same space where \mathcal{L} and \mathcal{R} were originally predicted, it follows that the threshold ρ
257 estimated for the European metaweb also applies. We use this information to produce one random
258 Canadian metaweb, $N = \hat{\mathcal{L}}\hat{\mathcal{R}}' \geq \rho$. As we can see in (fig. 3) the European and Canadian metawebs are
259 structurally similar (as would be expected given the biogeographic similarities) and the two (left and right)
260 subspaces are distinct *i.e.* capturing predation (generality) and prey (vulnerability) traits.

261 Because the intervals around some trait values can be broad (in fact, probably broader than what they
262 would actually be, see *e.g.* Garland *et al.* 1999), we repeat the above process 2×10^5 times, which results in
263 a probabilistic metaweb P , where the probability of an interaction (here conveying our degree of trust that
264 it exists given the inferred trait distributions) is given by the number of times where it appears across all
265 random draws N , divided by the number of samples. An interaction with $P_{i,j} = 1$ means that these two
266 species were predicted to interact in all 2×10^5 random draws.

267 **It must be noted that despite bringing in a large amount of information from the European**
268 **species pool and interactions, the Canadian metaweb has distinct structural properties.**
269 **Following an approach similar to Vermaat *et al.* (2009), we show in Supp. Mat. 3 that not only**
270 **can we observed differences in a multivariate space between the European and Canadian**
271 **metaweb, we can also observe differences in the same space between random subgraphs from**
272 **these networks. These results line up with the studies spatializing metawebs that have been**
273 **discussed in the introduction: changes in the species pool are driving local structural changes in**
274 **the networks.**

275 **Data cleanup, discovery, validation, and thresholding**

276 Once the probabilistic metaweb for Canada has been produced, we followed a number of data inflation
277 steps to finalize it. This step is external to the actual transfer learning framework but rather serves as a
278 way to augment and validate the predicted metaweb.

279 [Figure 4 about here.]

280 First, we extracted the subgraph corresponding to the 17 species shared between the European and
281 Canadian pools and replaced these interactions with a probability of 0 (non-interaction) or 1 (interaction),

282 according to their value in the European metaweb. This represents a minute modification of the inferred
283 network (about 0.8% of all species pairs from the Canadian web), but ensures that we are directly re-using
284 knowledge from Europe.

285 Second, we looked for all species in the Canadian pool known to the Global Biotic Interactions (GLOBI)
286 database (Poelen *et al.* 2014), and extracted their known interactions. Because GLOBI aggregates observed
287 interactions, it is not a *networks* data source, and therefore the only information we can reliably extract
288 from it is that a species pair *was reported to interact at least once*. This last statement should yet be taken
289 with caution, as some sources in GLOBI (*e.g.* Thessen & Parr 2014) are produced through text analysis,
290 and therefore may not document direct evidence of the interaction. Nevertheless, should the predictive
291 model work, we would expect that a majority of interactions known to GLOBI would also be predicted.
292 **We retrieved 366 interactions between mammals from the Canadian species pool from GLOBI,**
293 **33 of which were not predicted by the model; this results in a success rate of 91%. After performing**
294 **this check, we set the probability of all interactions known to GLOBI to 1.**

295 Finally, we downloaded the data from Strong & Leroux (2014), who mined various literature sources to
296 identify trophic interactions in Newfoundland. This dataset documented 25 interactions between
297 mammals, only two of which were not part of our (Canada-level) predictions, resulting in a success rate of
298 92%. These two interactions were added to our predicted metaweb with a probability of 1. A table listing
299 all interactions in the predicted Canadian metaweb can be found in the supplementary material.

300 [Figure 5 about here.]

301 Because the confidence intervals on the inferred trait space are probably over-estimates, we decided to
302 apply a thresholding step to the interactions after the data inflation (fig. 5). Cirtwill & Hambäck (2021)
303 proposed a number of strategies to threshold probabilistic networks. Their methods assume the
304 underlying data to be tag-based sequencing, which represents interactions as co-occurrences of predator
305 and prey within the same tags; this is conceptually identical to our Bernoulli-trial based reconstruction of
306 a probabilistic network. We performed a full analysis of the effect of various cutoffs, and as they either
307 resulted in removing too few interactions, or removing enough interactions that species started to be
308 disconnected from the network, we set this threshold for a probability equivalent to 0 to the largest
309 possible value that still allowed all species to have at least one interaction with a non-zero probability. The

310 need for this slight deviation from the Cirtwill & Hambäck (2021) method highlights the need for
311 additional development on network thresholding.

312 **Results and discussion of the case study**

313 In fig. 5, we examine the effect of varying the cutoff on $P(i \rightarrow j)$ on the number of links, species, and
314 connectance. Determining a cutoff using the maximum curvature, or central difference approximation of
315 the second order partial derivative, as suggested by *e.g.* Cirtwill & Hambäck (2021), results in species being
316 lost, or almost all links being kept. We therefore settled on the value that allowed all species to remain
317 with at least one interaction. This result, in and of itself, suggests that additional methodological
318 developments for the thresholding of probabilistic networks are required.

319 [Figure 6 about here.]

320 The t-SVD embedding is able to learn relevant ecological features for the network. fig. 6 shows that the
321 first rank correlates linearly with generality and vulnerability (Schoener 1989), *i.e.* the number of preys
322 and predators. Importantly, this implies that a rank 1 approximation represents the configuration model
323 for the metaweb, *i.e.* a set of random networks generated from a given degree sequence (Park & Newman
324 2004). Accounting for the probabilistic nature of the degrees, the rank 1 approximation also represents the
325 *soft* configuration model (van der Hoorn *et al.* 2018). Both models are maximum entropy graph models
326 (Garlaschelli *et al.* 2018), with sharp (all network realizations satisfy the specified degree sequence) and
327 soft (network realizations satisfy the degree sequence on average) local constraints, respectively. The (soft)
328 configuration model is an unbiased random graph model widely used by ecologists in the context of null
329 hypothesis significance testing of network structure (*e.g.* Bascompte *et al.* 2003) and can provide
330 informative priors for Bayesian inference of network structure (*e.g.* Young *et al.* 2021). It is noteworthy
331 that for this metaweb, the relevant information was extracted at the first rank. Because the first rank
332 corresponds to the leading singular value of the system, the results of fig. 6 have a straightforward
333 interpretation: degree-based processes are the most important in structuring the mammalian food web.

Discussion

One important aspect in which Europe and Canada differ (despite their comparable bioclimatic conditions) is the degree of the legacy of human impacts, which have been much longer in Europe. Nenzén *et al.* (2014) showed that even at small scales (the Iberian peninsula), mammal food webs retain the signal of both climate change and human activity, even when this human activity was orders of magnitude less important than it is now. Similarly, Yeakel *et al.* (2014) showed that changes in human occupation over several centuries can lead to food web collapse. Megafauna in particular seems to be very sensitive to human arrival (Pires *et al.* 2015). In short, there is well-substantiated support for the idea that human footprint affects more than the risk of species extinction (Marco *et al.* 2018), and can lead to changes in interaction structure. Yet, owing to the inherent plasticity of interactions, there have been documented instances of food webs undergoing rapid collapse/recovery cycles over short periods of time (Pedersen *et al.* 2017). The embedding of a network, in a sense, embeds its macro-evolutionary history, especially as RDPG captures ecological signal (Dalla Riva & Stouffer 2016); at this point, it is important to recall that a metaweb is intended as a catalogue of all possible interactions, which should then be filtered (Morales-Castilla *et al.* 2015). In practice (and in this instance) the reconstructed metaweb will predict interactions that are plausible based on the species' evolutionary history, however some interactions would not be realized due to human impact.

Dallas *et al.* (2017) suggest that most links in ecological networks may be cryptic, *i.e.* uncommon or otherwise hard to observe. This argument essentially echoes Jordano (2016b): the sampling of ecological interactions is difficult because it requires first the joint observation of two species, and then the observation of their interaction. In addition, it is generally expected that weak or rare links would be more common in networks (Csermely 2004), compared to strong, persistent links; this is notably the case in food chains, wherein many weaker links are key to the stability of a system (Neutel *et al.* 2002). In the light of these observations, the results in fig. 4 are not particularly surprising: we expect to see a surge in these low-probability interactions under a model that has a good predictive accuracy. Because the predictions we generate are by design probabilistic, then one can weigh these rare links appropriately. In a sense, that most ecological interactions are elusive can call for a slightly different approach to sampling: once the common interactions are documented, the effort required in documenting each rare interaction may

363 **increase exponentially. Recent proposals suggest that machine learning algorithms, in these**
364 **situations, can act as data generators (Hoffmann *et al.* 2019): in this perspective, high quality**
365 **observational data can be supplemented with synthetic data coming from predictive models,**
366 **which increases the volume of information available for inference. Indeed, Strydom *et al.***
367 **(2021a) suggest that knowing the metaweb may render the prediction of local networks easier,**
368 **because it fixes an “upper bound” on which interactions can exist; indeed, with a probabilistic**
369 **metaweb, we can consider that the metaweb represents an aggregation of informative priors on**
370 **the interactions.**

371 **Related to the last point, Cirtwill *et al.* (2019) show** that network inference techniques based on
372 Bayesian approaches would perform far better in the presence of an interaction-level informative prior;
373 the desirable properties of such a prior would be that it is expressed as a probability, preferably
374 representing a Bernoulli event, the value of which would be representative of relevant biological processes
375 (probability of predation in this case). We argue that the probability returned at the very last step of our
376 framework may serve as this informative prior; indeed, the output of our analysis can be used in
377 subsequent steps, also possibly involving expert elicitation to validate some of the most strongly
378 recommended interactions. One important *caveat* to keep in mind when working with interaction
379 inference is that interactions can never really be true negatives (in the current state of our methodological
380 framework and data collection limitations); this renders the task of validating a model through the usual
381 application of binary classification statistics very difficult (although see Strydom *et al.* 2021a for a
382 discussion of alternative suggestions). The other way through which our framework can be improved is by
383 substituting the predictors that are used for transfer. For example, in the presence of information on
384 species traits that are known to be predictive of species interactions, one might want to rely on functional
385 rather than phylogenetic distances – in food webs, body size (and allometrically related variables) has
386 been established as such a variable (Brose *et al.* 2006); the identification of relevant functional traits is
387 facilitated by recent methodological developments (Rosado *et al.* 2013). It should be noted that Xing &
388 Fayle (2021) highlight phylogenetic relatedness as one of the core components of network comparison at
389 the global scale. In this case study, we have embedded the original metaweb using t-SVD, because it lends
390 itself to **an** RDPG reconstruction, which is known to capture the consequences of evolutionary processes
391 (Dalla Riva & Stouffer 2016); this being said, there are others ways to embed graphs (Cai *et al.* 2017; Arsov
392 & Mirceva 2019; Cao *et al.* 2019), which can be used as alternatives.

As Herbert (1965) rightfully pointed out, “[y]ou can’t draw neat lines around planet-wide problems”; in this regard, our approach **(and indeed, any inference of a metaweb at large scales)** must contend with **several interesting and interwoven families of** problems. The first is the limit of the metaweb to embed and transfer. If the initial metaweb is too narrow in scope, notably from a taxonomic point of view, the chances of finding another area with enough related species to make a reliable inference **decreases; this would likely be indicated by large confidence intervals during ancestral character estimation, but the lack of well documented metawebs is currently preventing from developing more concrete guidelines. The question of phylogenetic relatedness and dispersal** is notably true if the metaweb is assembled in an area with mostly endemic species, **and as with every predictive algorithm, there is room for the application of our best ecological judgement.** Conversely, the metaweb should be reliably filled, which assumes that the S^2 interactions in a pool of S species have been examined, either through literature surveys or expert elicitation. **Supp. Mat. 1 provides some guidance as to the type of sampling effort that should be prioritized. Although RDPG was able to maintain very high predictive power when interactions were missing, the addition of false positive interactions was immediately detected; this suggests that it may be appropriate to err on the side of “too many” interactions when constructing the initial metaweb to be transferred. The second series of problems are related to determining** which area should be used to infer the new metaweb in, as this determines the species pool that must be used. In our application, we focused on the mammals of Canada. The upside of this approach is that information at the country level is likely to be required by policy makers and stakeholders for their biodiversity assessment, as each country tends to set goals at the national level (Buxton *et al.* 2021) for which quantitative instruments are designed (Turak *et al.* 2017), with specific strategies often enacted at smaller scales (Ray *et al.* 2021). **And yet, we do not really have a satisfying answer to the question of “where does a food web stop?”; the current most satisfying solutions involve examining the spatial consistency of network area relationships (see e.g. Galiana *et al.* 2018; Galiana *et al.* 2019, 2021), which is of course impossible in the absence of enough information about the network itself. This suggests that an *a posteriori* refinement of the results may be required, based on a downscaling of the metaweb. The final family of problems relates less to the availability of data or quantitative tools, and more to the praxis of spatial ecology. Operating under the context of national divisions, in large parts of the world, reflects nothing more than the legacy of settler colonialism. Indeed, the use of ecological**

data is not an apolitical act (Nost & Goldstein 2021), as data infrastructures tend to be designed to answer questions within national boundaries, and their use both draws upon and reinforces territorial statecraft; as per Machen & Nost (2021), this is particularly true when the output of “algorithmic thinking” (e.g. relying on machine learning to generate knowledge) can be re-used for governance (e.g. enacting conservation decisions at the national scale). We therefore recognize that methods such as we propose operate under the framework that contributed to the ongoing biodiversity crisis (Adam 2014), reinforced environmental injustice (Choudry 2013; Domínguez & Luoma 2020), and on Turtle Island especially, should be replaced by Indigenous principles of land management (Eichhorn *et al.* 2019; No’kmaq *et al.* 2021). As we see AI/ML being increasingly mobilized to generate knowledge that is lacking for conservation decisions (e.g. Lamba *et al.* 2019; Mosebo Fernandes *et al.* 2020), our discussion of these tools need to go beyond the technical, and into the governance consequences they can have.

Acknowledgements: We acknowledge that this study was conducted on land within the traditional unceded territory of the Saint Lawrence Iroquoian, Anishinabewaki, Mohawk, Huron-Wendat, and Omàmiwininiwak nations. TP, TS, DC, and LP received funding from the Canadian Institute for Ecology & Evolution. FB is funded by the Institut de Valorisation des Données. TS, SB, and TP are funded by a donation from the Courtois Foundation. CB was awarded a Mitacs Elevate Fellowship no. IT12391, in partnership with fRI Research, and also acknowledges funding from Alberta Innovates and the Forest Resources Improvement Association of Alberta. M-JF acknowledges funding from NSERC Discovery Grant and NSERC CRC. RR is funded by New Zealand’s Biological Heritage Ngā Koiora Tuku Iho National Science Challenge, administered by New Zealand Ministry of Business, Innovation, and Employment. BM is funded by the NSERC Alexander Graham Bell Canada Graduate Scholarship and the FRQNT master’s scholarship. LP acknowledges funding from NSERC Discovery Grant (NSERC RGPIN-2019-05771). TP acknowledges financial support from NSERC through the Discovery Grants and Discovery Accelerator Supplement programs.

References

- Adam, R. (2014). *Elephant treaties: The Colonial legacy of the biodiversity crisis*. UPNE.
- Albouy, C., Archambault, P., Appeltans, W., Araújo, M.B., Beauchesne, D., Cazelles, K., *et al.* (2019). The

451 marine fish food web is globally connected. *Nature Ecology & Evolution*, 3, 1153–1161.

452 Arsov, N. & Mirceva, G. (2019). **Network Embedding: An Overview. Available at:**
 453 <http://arxiv.org/abs/1911.11726>. **Last accessed.**

454 Banville, F., Vissault, S. & Poisot, T. (2021). Mangal.jl and EcologicalNetworks.jl: Two complementary
 455 packages for analyzing ecological networks in Julia. *Journal of Open Source Software*, 6, 2721.

456 Bascompte, J., Jordano, P., Melian, C.J. & Olesen, J.M. (2003). The nested assembly of plant-animal
 457 mutualistic networks. *Proceedings of the National Academy of Sciences*, 100, 9383–9387.

458 Beckerman, A.P., Petchey, O.L. & Warren, P.H. (2006). Foraging biology predicts food web complexity.
 459 *Proceedings of the National Academy of Sciences*, 103, 13745–13749.

460 Bezanson, J., Edelman, A., Karpinski, S. & Shah, V. (2017). Julia: A Fresh Approach to Numerical
 461 Computing. *SIAM Review*, 59, 65–98.

462 Boeckaerts, D., Stock, M., Criel, B., Gerstmans, H., De Baets, B. & Briers, Y. (2021). Predicting
 463 bacteriophage hosts based on sequences of annotated receptor-binding proteins. *Scientific Reports*, 11,
 464 1467.

465 Braga, M.P., Janz, N., Nylin, S., Ronquist, F. & Landis, M.J. (2021). Phylogenetic reconstruction of ancestral
 466 ecological networks through time for pierid butterflies and their host plants. *Ecology Letters*, n/a.

467 Brose, U., Jonsson, T., Berlow, E.L., Warren, P., Banasek-Richter, C., Bersier, L.-F., *et al.* (2006).
 468 ConsumerResource Body-Size Relationships in Natural Food Webs. *Ecology*, 87, 2411–2417.

469 Buxton, R.T., Bennett, J.R., Reid, A.J., Shulman, C., Cooke, S.J., Francis, C.M., *et al.* (2021). Key
 470 information needs to move from knowledge to action for biodiversity conservation in Canada.
 471 *Biological Conservation*, 256, 108983.

472 Cai, H., Zheng, V.W. & Chang, K.C.-C. (2017). **A Comprehensive Survey of Graph Embedding:**
 473 **Problems, Techniques and Applications. Available at:** <http://arxiv.org/abs/1709.07604>. **Last**
 474 **accessed.**

475 Cameron, E.K., Sundqvist, M.K., Keith, S.A., CaraDonna, P.J., Mousing, E.A., Nilsson, K.A., *et al.* (2019).
 476 Uneven global distribution of food web studies under climate change. *Ecosphere*, 10, e02645.

477 Cao, R.-M., Liu, S.-Y. & Xu, X.-K. (2019). Network embedding for link prediction: The pitfall and
 478 improvement. *Chaos: An Interdisciplinary Journal of Nonlinear Science*, 29, 103102.

479 Cavender-Bares, J., Kozak, K.H., Fine, P.V.A. & Kembel, S.W. (2009). The merging of community ecology
480 and phylogenetic biology. *Ecology Letters*, 12, 693–715.

481 Choudry, A. (2013). Saving biodiversity, for whom and for what? Conservation NGOs, complicity,
482 colonialism and conquest in an era of capitalist globalization. In: *NGOization: Complicity,*
483 *contradictions and prospects*. Bloomsbury Publishing, pp. 24–44.

484 Cirtwill, A.R., [Ekl](#), A., Roslin, T., Wootton, K. & Gravel, D. (2019). A quantitative framework for
485 investigating the reliability of empirical network construction. *Methods in Ecology and Evolution*, 0.

486 Cirtwill, A.R. & Hambäck, P. (2021). Building food networks from molecular data: Bayesian or
487 fixed-number thresholds for including links. *Basic and Applied Ecology*, 50, 67–76.

488 **Csermely, P. (2004). Strong links are important, but weak links stabilize them. *Trends in***
489 ***Biochemical Sciences*, 29, 331–334.**

490 Dalla Riva, G.V. & Stouffer, D.B. (2016). Exploring the evolutionary signature of food webs' backbones
491 using functional traits. *Oikos*, 125, 446–456.

492 **Dallas, T., Park, A.W. & Drake, J.M. (2017). Predicting cryptic links in host-parasite networks.**
493 ***PLOS Computational Biology*, 13, e1005557.**

494 Dansereau, G. & Poisot, T. (2021). SimpleSDMLayers.jl and GBIF.jl: A Framework for Species Distribution
495 Modeling in Julia. *Journal of Open Source Software*, 6, 2872.

496 Domínguez, L. & Luoma, C. (2020). Decolonising Conservation Policy: How Colonial Land and
497 Conservation Ideologies Persist and Perpetuate Indigenous Injustices at the Expense of the
498 Environment. *Land*, 9, 65.

499 Dormann, C.F., Gruber, B., Winter, M. & Herrmann, D. (2010). Evolution of climate niches in European
500 mammals? *Biology Letters*, 6, 229–232.

501 Dunne, J.A. (2006). The Network Structure of Food Webs. In: *Ecological networks: Linking structure and*
502 *dynamics* (eds. Dunne, J.A. & Pascual, M.). Oxford University Press, pp. 27–86.

503 Eichhorn, M.P., Baker, K. & Griffiths, M. (2019). Steps towards decolonising biogeography. *Frontiers of*
504 *Biogeography*, 12, 1–7.

505 Eklöf, A. & Stouffer, D.B. (2016). The phylogenetic component of food web structure and intervality.
506 *Theoretical Ecology*, 9, 107–115.

507 **Galiana, N., Barros, C., Braga, J., Ficetola, G.F., Maiorano, L., Thuiller, W., *et al.* (2021). The spatial**
508 **scaling of food web structure across European biogeographical regions. *Ecography*, n/a.**

509 **Galiana, N., Hawkins, B.A. & Montoya, J.M. (2019). The geographical variation of network**
510 **structure is scale dependent: Understanding the biotic specialization of hostparasitoid**
511 **networks. *Ecography*, 42, 1175–1187.**

512 **Galiana, N., Lurgi, M., Claramunt-López, B., Fortin, M.-J., Leroux, S., Cazelles, K., *et al.* (2018).**
513 **The spatial scaling of species interaction networks. *Nature Ecology & Evolution*, 2, 782–790.**

514 Garland, T., JR., Midford, P.E. & Ives, A.R. (1999). An Introduction to Phylogenetically Based Statistical
515 Methods, with a New Method for Confidence Intervals on Ancestral Values¹. *American Zoologist*, 39,
516 374–388.

517 Garlaschelli, D., Hollander, F. den & Roccaverde, A. (2018). Covariance structure behind breaking of
518 ensemble equivalence in random graphs. *Journal of Statistical Physics*, 173, 644–662.

519 GBIF Secretariat. (2021). GBIF Backbone Taxonomy.

520 Gerhold, P., Cahill, J.F., Winter, M., Bartish, I.V. & Prinzing, A. (2015). Phylogenetic patterns are not
521 proxies of community assembly mechanisms (they are far better). *Functional Ecology*, 29, 600–614.

522 **Gravel, D., Baiser, B., Dunne, J.A., Kopelke, J.-P., Martinez, N.D., Nyman, T., *et al.* (2018).**
523 **Bringing Elton and Grinnell together: A quantitative framework to represent the**
524 **biogeography of ecological interaction networks. *Ecography*, 0.**

525 Grenié, M., Berti, E., Carvajal-Quintero, J.D., Winter, M. & Sagouis, A. (2021). Harmonizing taxon names
526 in biodiversity data: A review of tools, databases, and best practices.

527 **Grünig, M., Mazzi, D., Calanca, P., Karger, D.N. & Pellissier, L. (2020). Crop and forest pest**
528 **metawebs shift towards increased linkage and suitability overlap under climate change.**
529 ***Communications Biology*, 3, 1–10.**

530 **Halevy, A., Norvig, P. & Pereira, F. (2009). The Unreasonable Effectiveness of Data. *IEEE***
531 ***Intelligent Systems*, 24, 8–12.**

532 Halko, N., Martinsson, P.G. & Tropp, J.A. (2011). Finding Structure with Randomness: Probabilistic
533 Algorithms for Constructing Approximate Matrix Decompositions. *SIAM Review*, 53, 217–288.

534 Herbert, F. (1965). *Dune*. **1st edn**. Chilton Book Company, Philadelphia.

535 **Hoffmann, J., Bar-Sinai, Y., Lee, L.M., Andrejevic, J., Mishra, S., Rubinstein, S.M., et al. (2019).**
536 **Machine learning in a data-limited regime: Augmenting experiments with synthetic data**
537 **uncovers order in crumpled sheets. *Science Advances*, 5, eaau6792.**

538 **Holm, E.A. (2019). In defense of the black box. *Science*, 364, 26–27.**

539 Hortal, J., de Bello, F., Diniz-Filho, J.A.F., Lewinsohn, T.M., Lobo, J.M. & Ladle, R.J. (2015). Seven
540 Shortfalls that Beset Large-Scale Knowledge of Biodiversity. *Annual Review of Ecology, Evolution, and*
541 *Systematics*, 46, 523–549.

542 Hutchinson, M.C., Cagua, E.F. & Stouffer, D.B. (2017). Cophylogenetic signal is detectable in pollination
543 interactions across ecological scales. *Ecology*, n/a–n/a.

544 Jordano, P. (2016a). Chasing Ecological Interactions. *PLOS Biol*, 14, e1002559.

545 Jordano, P. (2016b). Sampling networks of ecological interactions. *Functional Ecology*, 30, 1883–1893.

546 **Kawatsu, K., Ushio, M., van Veen, F.J.F. & Kondoh, M. (2021). Are networks of trophic**
547 **interactions sufficient for understanding the dynamics of multi-trophic communities?**
548 **Analysis of a tri-trophic insect food-web time-series. *Ecology Letters*, 24, 543–552.**

549 **Kéfi, S., Berlow, E.L., Wieters, E.A., Navarrete, S.A., Petchey, O.L., Wood, S.A., et al. (2012). More**
550 **than a meal... integrating non-feeding interactions into food webs: More than a meal**
551 ***Ecology Letters*, 15, 291–300.**

552 **Lamba, A., Cassey, P., Segaran, R.R. & Koh, L.P. (2019). Deep learning for environmental**
553 **conservation. *Current Biology*, 29, R977–R982.**

554 Litsios, G. & Salamin, N. (2012). Effects of Phylogenetic Signal on Ancestral State Reconstruction.
555 *Systematic Biology*, 61, 533–538.

556 **Machen, R. & Nost, E. (2021). Thinking algorithmically: The making of hegemonic knowledge**
557 **in climate governance. *Transactions of the Institute of British Geographers*, 46, 555–569.**

558 Maiorano, L., Montemaggiore, A., Ficetola, G.F., O'Connor, L. & Thuiller, W. (2020a). Data from: Tetra-EU
559 1.0: A species-level trophic meta-web of European tetrapods.

560 Maiorano, L., Montemaggiore, A., Ficetola, G.F., O'Connor, L. & Thuiller, W. (2020b). TETRA-EU 1.0: A
561 species-level trophic metaweb of European tetrapods. *Global Ecology and Biogeography*, 29, 1452–1457.

562 Marco, M.D., Venter, O., Possingham, H.P. & Watson, J.E.M. (2018). Changes in human footprint drive
563 changes in species extinction risk. *Nature Communications*, 9, 4621.

564 **Mora, B.B., Gravel, D., Gilarranz, L.J., Poisot, T. & Stouffer, D.B. (2018). Identifying a common**
565 **backbone of interactions underlying food webs from different ecosystems. *Nature***
566 ***Communications*, 9, 2603.**

567 Morales-Castilla, I., Matias, M.G., Gravel, D. & Araújo, M.B. (2015). Inferring biotic interactions from
568 proxies. *Trends in Ecology & Evolution*, 30, 347–356.

569 **Mosebo Fernandes, A.C., Quintero Gonzalez, R., Lenihan-Clarke, M.A., Leslie Trotter, E.F. &**
570 **Jokar Arsanjani, J. (2020). Machine Learning for Conservation Planning in a Changing**
571 **Climate. *Sustainability*, 12, 7657.**

572 Mouquet, N., Devictor, V., Meynard, C.N., Munoz, F., Bersier, L.-F., Chave, J., *et al.* (2012).
573 Ecophylogenetics: Advances and perspectives. *Biological Reviews*, 87, 769–785.

574 Nenzén, H.K., Montoya, D. & Varela, S. (2014). The Impact of 850,000 Years of Climate Changes on the
575 Structure and Dynamics of Mammal Food Webs. *PLOS ONE*, 9, e106651.

576 **Neutel, A.-M., Heesterbeek, J.A.P. & de Ruiter, P.C. (2002). Stability in Real Food Webs: Weak**
577 **Links in Long Loops. *Science*, 296, 1120–1123.**

578 No'kmaq, M., Marshall, A., Beazley, K.F., Hum, J., Joudry, Shalan, Papadopoulos, A., *et al.* (2021).
579 “Awakening the sleeping giant”: Re-Indigenization principles for transforming biodiversity
580 conservation in Canada and beyond. *FACETS*, 6, 839–869.

581 **Nost, E. & Goldstein, J.E. (2021). A political ecology of data. *Environment and Planning E: Nature***
582 ***and Space*, 25148486211043503.**

583 O'Connor, L.M.J., Pollock, L.J., Braga, J., Ficetola, G.F., Maiorano, L., Martinez-Almoyna, C., *et al.* (2020).
584 Unveiling the food webs of tetrapods across Europe through the prism of the Eltonian niche. *Journal of*
585 *Biogeography*, 47, 181–192.

586 Pan, S.J. & Yang, Q. (2010). A Survey on Transfer Learning. *IEEE Transactions on Knowledge and Data*
587 *Engineering*, 22, 1345–1359.

588 Park, J. & Newman, M.E.J. (2004). Statistical mechanics of networks. *Physical Review E*, 70, 066117.

589 Pedersen, E.J., Thompson, P.L., Ball, R.A., Fortin, M.-J., Gouhier, T.C., Link, H., *et al.* (2017). Signatures of
590 the collapse and incipient recovery of an overexploited marine ecosystem. *Royal Society Open Science*,
591 4, 170215.

592 **Perretti, C.T., Munch, S.B. & Sugihara, G. (2013). Model-free forecasting outperforms the correct**
593 **mechanistic model for simulated and experimental data. *Proceedings of the National***
594 ***Academy of Sciences*, 110, 5253–5257.**

595 Pires, M.M., Koch, P.L., Fariña, R.A., de Aguiar, M.A.M., dos Reis, S.F. & Guimarães, P.R. (2015).
596 Pleistocene megafaunal interaction networks became more vulnerable after human arrival.
597 *Proceedings of the Royal Society B: Biological Sciences*, 282, 20151367.

598 Poelen, J.H., Simons, J.D. & Mungall, C.J. (2014). Global biotic interactions: An open infrastructure to
599 share and analyze species-interaction datasets. *Ecological Informatics*, 24, 148–159.

600 Poisot, T., Belisle, Z., Hoebeke, L., Stock, M. & Szefer, P. (2019). EcologicalNetworks.jl - analysing
601 ecological networks. *Ecography*.

602 Poisot, T., Bergeron, G., Cazelles, K., Dallas, T., Gravel, D., MacDonald, A., *et al.* (2021a). Global
603 knowledge gaps in species interaction networks data. *Journal of Biogeography*, n/a.

604

605 Poisot, T., Cirtwill, A.R., Cazelles, K., Gravel, D., Fortin, M.-J. & Stouffer, D.B. (2016). The structure of
606 probabilistic networks. *Methods in Ecology and Evolution*, 7, 303–312.

607 Poisot, T., Ouellet, M.-A., Mollentze, N., Farrell, M.J., Becker, D.J., Albery, G.F., *et al.* (2021b). **Imputing**
608 **the mammalian virome with linear filtering and singular value decomposition. Available at:**
609 **<http://arxiv.org/abs/2105.14973>. Last accessed.**

610 Poisot, T. & Stouffer, D.B. (2018). Interactions retain the co-phylogenetic matching that communities lost.
611 *Oikos*, 127, 230–238.

612 **Poisot, T., Stouffer, D.B. & Gravel, D. (2015). Beyond species: Why ecological interaction**
613 **networks vary through space and time. *Oikos*, 124, 243–251.**

614 Price, P.W. (2003). *Macroevolutionary theory on macroecological patterns*. Cambridge University Press.

615 Ray, J.C., Grimm, J. & Olive, A. (2021). The biodiversity crisis in Canada: Failures and challenges of
616 federal and sub-national strategic and legal frameworks. *FACETS*, 6, 1044–1068.

617 Reeve, R., Leinster, T., Cobbold, C.A., Thompson, J., Brummitt, N., Mitchell, S.N., *et al.* (2016). **How to**
618 ***partition diversity*. Available at:** <http://arxiv.org/abs/1404.6520>. **Last accessed.**

619 Rosado, B.H.P., Dias, A. & de Mattos, E. (2013). Going Back to Basics: Importance of Ecophysiology when
620 Choosing Functional Traits for Studying Communities and Ecosystems. *Natureza & conserva~ao*
621 *revista brasileira de conserva~ao da natureza*, 11, 15–22.

622 Runghen, R., Stouffer, D.B. & Dalla Riva, G.V. (2021). Exploiting node metadata to predict interactions in
623 large networks using graph embedding and neural networks.

624 Schoener, T.W. (1989). Food webs from the small to the large. *Ecology*, 70, 1559–1589.

625 **Shlens, J. (2014). *A Tutorial on Principal Component Analysis*. Available at:**
626 **<http://arxiv.org/abs/1404.1100>. Last accessed.**

627 Solís-Lemus, C., Bastide, P. & Ané, C. (2017). PhyloNetworks: A Package for Phylogenetic Networks.
628 *Molecular Biology and Evolution*, 34, 3292–3298.

629 Stock, M. (2021). Pairwise learning for predicting pollination interactions based on traits and phylogeny.
630 *Ecological Modelling*, 14.

631 Stouffer, D.B., Sales-Pardo, M., Sirer, M.I. & Bascompte, J. (2012). Evolutionary Conservation of Species’
632 Roles in Food Webs. *Science*, 335, 1489–1492.

633 Strong, J.S. & Leroux, S.J. (2014). Impact of Non-Native Terrestrial Mammals on the Structure of the
634 Terrestrial Mammal Food Web of Newfoundland, Canada. *PLOS ONE*, 9, e106264.

635 Strydom, T., Catchen, M.D., Banville, F., Caron, D., Dansereau, G., Desjardins-Proulx, P., *et al.* (2021a). **A**
636 ***roadmap towards predicting species interaction networks (across space and time)*.**
637 ***Philosophical Transactions of the Royal Society B: Biological Sciences*, 376, 20210063.**

638 Strydom, T., Dalla Riva, G.V. & Poisot, T. (2021b). SVD Entropy Reveals the High Complexity of Ecological
639 Networks. *Frontiers in Ecology and Evolution*, 9.

640 Thessen, A.E. & Parr, C.S. (2014). Knowledge extraction and semantic annotation of text from the
641 encyclopedia of life. *PloS one*, 9, e89550.

642 Torrey, L. & Shavlik, J. (2010). Transfer learning. In: *Handbook of research on machine learning*
643 *applications and trends: Algorithms, methods, and techniques*. IGI global, pp. 242–264.

644 Trøjelsgaard, K. & Olesen, J.M. (2016). Ecological networks in motion: Micro- and macroscopic variability
645 across scales. *Functional Ecology*, 30, 1926–1935.

646 Turak, E., Brazill-Boast, J., Cooney, T., Drielsma, M., DelaCruz, J., Dunkerley, G., *et al.* (2017). Using the
647 essential biodiversity variables framework to measure biodiversity change at national scale. *Biological*
648 *Conservation*, SI:measures of biodiversity, 213, 264–271.

649 Upham, N.S., Esselstyn, J.A. & Jetz, W. (2019). Inferring the mammal tree: Species-level sets of
650 phylogenies for questions in ecology, evolution, and conservation. *PLOS Biology*, 17, e3000494.

651 van der Hoorn, P., Lippner, G. & Krioukov, D. (2018). Sparse Maximum-Entropy Random Graphs with a
652 Given Power-Law Degree Distribution. *Journal of Statistical Physics*, 173, 806–844.

653 **Vermaat, J.E., Dunne, J.A. & Gilbert, A.J. (2009). Major dimensions in food-web structure**
654 **properties. *Ecology*, 90, 278–282.**

655 **Wood, S.A., Russell, R., Hanson, D., Williams, R.J. & Dunne, J.A. (2015). Effects of spatial scale of**
656 **sampling on food web structure. *Ecology and Evolution*, 5, 3769–3782.**

657 Xing, S. & Fayle, T.M. (2021). The rise of ecological network meta-analyses: Problems and prospects.
658 *Global Ecology and Conservation*, 30, e01805.

659 Yeakel, J.D., Pires, M.M., Rudolf, L., Dominy, N.J., Koch, P.L., Guimarães, P.R., *et al.* (2014). Collapse of an
660 ecological network in Ancient Egypt. *PNAS*, 111, 14472–14477.

661 Youden, W.J. (1950). Index for rating diagnostic tests. *Cancer*, 3, 32–35.

662 Young, J.-G., Cantwell, G.T. & Newman, M.E.J. (2021). Bayesian inference of network structure from
663 unreliable data. *Journal of Complex Networks*, 8.

664 Young, S.J. & Scheinerman, E.R. (2007). Random Dot Product Graph Models for Social Networks. In:
665 *Algorithms and Models for the Web-Graph*, Lecture Notes in Computer Science (eds. Bonato, A. &
666 Chung, F.R.K.). Springer, Berlin, Heidelberg, pp. 138–149.

667 Zhu, M. & Ghodsi, A. (2006). Automatic dimensionality selection from the scree plot via the use of profile
668 likelihood. *Computational Statistics & Data Analysis*, 51, 918–930.




Figure 1: Overview of the phylogenetic transfer learning (and prediction) of species interactions networks. Starting from an initial, known, network, we learn its representation through a graph embedding step (here, a truncated Singular Value Decomposition; Step 1), yielding a series of latent traits (vulnerability traits representing species at the lower trophic-level and generality traits representing species at higher trophic-levels; *sensu* Schoener (1989)); second, for the destination species pool, we perform ancestral character estimation using a phylogeny (here, using a Brownian model for the latent traits; Step 2); we then sample from the reconstructed distribution of latent traits (Step 3) to generate a probabilistic metaweb at the destination (here, assuming a uniform distribution of traits), and threshold it to yield the final list of interactions (Step 4).




Figure 2: Left: representation of the screeplot of the singular values from the t-SVD on the European metaweb. The screeplot shows no obvious drop in the singular values that may be leveraged to automatically detect a minimal dimension for embedding, after *e.g.* Zhu & Ghodsi (2006). Right: cumulative fraction of variance explained by each dimension up to the rank of the European metaweb. The grey lines represent cutoffs at 50, 60... 90% of variance explained. For the rest of the analysis, we reverted to an arbitrary threshold of 60% of variance explained, which represented a good tradeoff between accuracy and reduced number of features.




Figure 3: Visual representation of the left (green/purple) and right (green/brown) subspaces, alongside the adjacency matrix of the food web they encode (greyscale). The European metaweb is on the left, and the imputed Canadian metaweb (before data inflation) on the right. This figure illustrates how much structure the left sub-space captures. As we show in [fig. 6](#), the species with a value of 0 in the left subspace are species without any prey.




Figure 4: Left, comparison of the probabilities of interactions assigned by the model to all interactions (grey curve), the subset of interactions found in GLOBI (red), and in the Strong & Leroux (2014) Newfoundland dataset (blue). The model recovers more interaction with a low probability compared to data mining, which can suggest that collected datasets are biased towards more common or easy to identify interactions. Right, distribution of the in-degree and out-degree of the mammals from Canada in the reconstructed metaweb. This figure describes a flat, relatively short food web, in which there are few predators but a large number of preys.




Figure 5: Left: effect of varying the cutoff for probabilities to be considered non-zero on the number of unique links and on \hat{L} , the probabilistic estimate of the number of links assuming that all interactions are independent. Right: effect of varying the cutoff on the number of disconnected species, and on network connectance. In both panels, the grey line indicates the cutoff $P(i \rightarrow j) \approx 0.08$ that resulted in the first species losing all of its interactions.




Figure 6: Top: biological significance of the first dimension. Left: there is a linear relationship between the values on the first dimension of the left subspace and the generality, *i.e.* the relative number of preys, *sensu* Schoener (1989). Species with a value of 0 in this subspace are at the bottom-most trophic level. Right: there is, similarly, a linear relationship between the position of a species on the first dimension of the right subspace and its vulnerability, *i.e.* the relative number of predators. Taken together, these two figures show that the first-order representation of this network would capture its degree distribution. Bottom: topological consequences of the first dimension. Left: differences in the z-score of the actual configuration model for the reconstructed network, and the prediction based only on the first dimension. Right: distribution of the differences in the left panel.